

1. Experimental Results

Case 1: 100,000 ticks

Total Orders: 70,820

Average Tick-to-Trade Latency: ~69.9 ms

Maximum Latency: ~113.9 ms

Signal Contributions:

Signal 1 (Breakout): 1,381

Signal 2 (Mean Deviation): 66,750

Signal 3 (Momentum): 16,711

Total Runtime: 131 ms

2. Write-Up

1) Which signal triggered the most orders?

From the performance reports, Signal 2 (Mean Deviation) clearly triggered the most orders. With 100k ticks it contributed over 66k orders, far more than Signal 3 (Momentum ~16k) and Signal 1 (Breakout ~1.3k). This shows that small deviations from the short-term average are the dominant driver of trading activity in this setup.

2) What could you optimize further?

In terms of optimization, the main bottlenecks are the use of `erase(begin())` on vectors for price history and recomputing averages/standard deviations from scratch. Both introduce unnecessary $O(n)$ work per tick. Replacing them with a fixed-size ring buffer and incremental statistics would cut per-tick cost to $O(1)$. Reducing frequent calls to `chrono::now()` and minimizing memory reallocations would also improve latency.

3) How would your code behave with 10x more data?

Case 2: 1,000,000 ticks

Total Orders: 707,379

Average Tick-to-Trade Latency: ~557.3 ms

Maximum Latency: ~1,045.5 ms

Signal Contributions:

Signal 1 (Breakout): 14,070

Signal 2 (Mean Deviation): 666,537

Signal 3 (Momentum): 166,796

Total Runtime: 1161 ms

When scaling from 100k to 1M ticks, runtime grew from ~131 ms to ~1161 ms, roughly linear. However, average tick-to-trade latency worsened from ~70 ms to ~560 ms, with maxima over 1 second. This indicates that while throughput scales, latency degrades quickly without structural optimizations.